

## Testat 3: GPU Programming

Semester week 10

**Submission deadline:** 11.05.2023 Rapperswil  
19.05.2023 St Gallen

Submission over Moodle

- Working in a team allowed (max. 2 people). Please put the names on the submission
- What needs to be submitted
  1. Submit your source code
  2. a pdf document reporting the required performance measurements and answers to the questions.

### Exercise 10:

Goal:

- Optimize parallel GPU programs with advanced techniques.
- Shared memory and synchronization in CUDA resp. OpenCL.

#### Task 1: Shared Memory und Synchronization

Implement tiled matrix multiplication using shared memory as discussed in the lecture.

```
__global__  
void matrixMultKernel(float *A, float *B, float *C) {  
    __shared__ float Asub[TILE_SIZE][TILE_SIZE];  
    __shared__ float Bsub[TILE_SIZE][TILE_SIZE];  
  
    int tx = threadIdx.x, ty = threadIdx.y;  
    int col = blockIdx.x * TILE_SIZE + tx;  
    int row = blockIdx.y * TILE_SIZE + ty;  
    int nofTiles = (A_COLS + TILE_SIZE - 1) / TILE_SIZE;  
    //your code  
}
```

1. Measure the performance varying the tile size. **Report your measurements and the best tile size. Is there a relationship between the tile size and the speedup? Explain.**
2. Now Swap the dimensions row and column.

```
int    row    =    blockIdx.x    *    TILE_SIZE    +    tx;  
int    col    =    blockIdx.y    *    TILE_SIZE    +    ty;
```

and compare the run time. Vary the tile size. **Please report the experiments you performed, your measurements, speedup, plots, and the best size of the tiles for this case.**

3. (Optional) Building on step2, Measure the performance varying the size of the matrix (smaller and bigger), the block size, and the tile size. **Report your measurements and the best tile size for each matrix. Has the matrix size or block size any influence on the speedup? Explain.**

Recap: With tiled matrix multiplication, each block of matrix C is calculated in several synchronized stages. For each stage, the partial areas of A and B are first loaded into the shared memory in order to reduce repeated accesses in A and B.

